

Start zookeeper server.

```
Problems theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0 x
[2024-05-05 08:06:01,092] INFO Configuring NIO connection handler with 10s sessionle
ss connection timeout, 1 selector thread(s), 4 worker threads, and 64 kB direct buff
ers. (org.apache.zookeeper.server.NIOServerCnxnFactory)
[2024-05-05 08:06:01,098] INFO binding to port 0.0.0.0/0.0.0.0:2181 (org.apache.zook
eeper.server.NIOServerCnxnFactory)
[2024-05-05 08:06:01,116] INFO zookeeper.snapshotSizeFactor = 0.33 (org.apache.zooke
eper.server.ZKDatabase)
[2024-05-05 08:06:01,117] INFO Reading snapshot /tmp/zookeeper/version-2/snapshot.0
(org.apache.zookeeper.server.persistence.FileSnap)
[2024-05-05 08:06:01,123] INFO Snapshotting: 0x0 to /tmp/zookeeper/version-2/snasho
t.0 (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2024-05-05 08:06:01,145] INFO PrepRequestProcessor (sid:0) started, reconfigEnabled
=false (org.apache.zookeeper.server.PrepRequestProcessor)
[2024-05-05 08:06:01,155] INFO Using checkIntervalMs=60000 maxPerMinute=10000 (org.a
pache.zookeeper.server.ContainerManager)
```

Start Kafka server

```
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ bin/kafka-server-start.sh config/server.pr
operties
JVMJ9VM007W Command-line option unrecognized: -Xlog:gc*:file=/home/project/kafka_2.12-2.8.0/bin/../logs/k
afkaServer-gc.log:time,tags:filecount=10,filesize=100M
[2024-05-05 09:16:10,750] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log4jContro
llerRegistration$)
[2024-05-05 09:16:11,088] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disable clie
nt-initiated TLS renegotiation (org.apache.zookeeper.common.X509Util)
[2024-05-05 09:16:11,183] INFO Registered signal handlers for TERM, INT, HUP (org.apache.kafka.common.util
ls.LoggingSignalHandler)
[2024-05-05 09:16:11,188] INFO starting (kafka.server.KafkaServer)
```

3. Create a Kakfa topic named `toll`

```
e/project/kafka_2.12-2.8.0  theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0 x

theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ bin/kafka-topics.sh --create --topic toll --bootstrap-server localhost:9092 --replication-factor 2
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ bin/kafka-topics.sh --create --topic toll --bootstrap-server localhost:9092 --replication-factor 2 --partitions 2
Error while executing topic command : Replication factor: 2 larger than available brokers: 1.
[2024-05-05 09:12:28,446] ERROR org.apache.kafka.common.errors.InvalidReplicationFactorException: Replication factor: 2 larger than available brokers: 1.
(kafka.admin.TopicCommand$)
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ bin/kafka-topics.sh --create --topic toll --bootstrap-server localhost:9092 --replication-factor 1 --partitions 2
Created topic toll.
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$
```

Open the `toll_traffic_generator.py` and set the topic to `toll`.

```
kafka_2.12-2.8.0 > toll_traffic_generator.py > ...
1  """
2  Top Traffic Simulator
3  """
4  from time import sleep, time, ctime
5  from random import random, randint, choice
6  from kafka import KafkaProducer
7  producer = KafkaProducer(bootstrap_servers='localhost:9092')
8
9  TOPIC = 'toll'
10
11  VEHICLE_TYPES = ("car", "car", "car", "car", "car", "car",
12                  "car", "car", "car", "truck", "truck", "truck",
13                  "truck", "van", "van")
14  for i in range(100000):

kafka_2.12-2.8.0  theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0 x

Length: 828 [text/x-python]
Saving to: 'toll_traffic_generator.py'

toll_traffic_generator.  0%[                               ]    0  --
toll_traffic_generator. 100%[=====>]    828  --
.-KB/s    in 0s

2024-05-05 09:17:27 (171 MB/s) - 'toll_traffic_generator.py' saved [828/828]

theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$
```

Run the `toll_traffic_generator.py`.

```
File "toll_traffic_generator.py", line 23, in <module>
    sleep(random() * 2)
KeyboardInterrupt
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ python3
toll_traffic_generator.py
A truck has passed by the toll plaza 4009 at Sun May 5 09:33:54 2024.
A car has passed by the toll plaza 4001 at Sun May 5 09:33:54 2024.
A car has passed by the toll plaza 4002 at Sun May 5 09:33:55 2024.
A car has passed by the toll plaza 4000 at Sun May 5 09:33:56 2024.
A van has passed by the toll plaza 4000 at Sun May 5 09:33:56 2024.
█
```

Download the `streaming_data_reader.py` from the url below using 'wget'.

```
kafka_2.12-2.8.0 > streaming_data_reader.py > ...
4 from datetime import datetime
5 from kafka import KafkaConsumer
6 import mysql.connector
7
8 TOPIC='toll'
9 DATABASE = 'tolldata'
10 USERNAME = 'root'
11 PASSWORD = 'MTQxMjQtaGFzaGFh'
12
13 print("Connecting to the database")
14 try:
15     connection = mysql.connector.connect(host='localhost'
16 except Exception:
17     print("Error connecting to the database")

kafka_2.12-2.8.0 theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0 x █

+-----+
1 row in set (0.00 sec)

mysql> exit
exit
^C
mysql> ^Z
[1]+  Stopped                  mysql --host=127.0.0.1 --port=3306 --user
=root --password=MTQxMjQtaGFzaGFh
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$
theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0$ █
```

6. Run the `streaming_data_reader.py`

```
4 from datetime import datetime
5 from kafka import KafkaConsumer
6 import mysql.connector
7
8 TOPIC='toll'
9 DATABASE = 'tolldata'
10 USERNAME = 'root'
11 PASSWORD = 'MTQxMjQtaGFzaGFh'
12
13 print("Connecting to the database")
14 try:
15     connection = mysql.connector.connect(host='localhost'
16 except Exception:
17     print("Error connecting to the database")
```

kafka_2.12-2.8.0 theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0 x

[1]+ Stopped mysql --host=127.0.0.1 --port=3306 --user=root --password=MTQxMjQtaGFzaGFh

theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0\$

theia@theiadocker-hashaammusht:/home/project/kafka_2.12-2.8.0\$ python3 streaming_data_reader.py

Connecting to the database

Connected to database

Connecting to Kafka

Connected to Kafka

Reading messages from the topic toll

List the top 10 rows in the table `livetolldata`.

```
streaming_data_reader.py x
4 from datetime import datetime
5 from kafka import KafkaConsumer
6 import mysql.connector
7
8 TOPIC='toll'
9 DATABASE = 'tolldata'
10 USERNAME = 'root'
11 PASSWORD = 'MTQxMjQtaGFzaGFh'
12
13 print("Connecting to the database")
14 try:
```

kafka_2.12-2.8.0 theia@theiadocker-hashaammusht: /home/project/kafka_2.12-2.8.0

```
Database changed
mysql> select * from livetolldata limit 10;
```

timestamp	vehicle_id	vehicle_type	toll_plaza_id
2024-05-05 10:43:50	4803979	car	4008
2024-05-05 10:43:51	9518063	car	4006
2024-05-05 10:43:52	4691408	car	4007
2024-05-05 10:43:52	8704613	truck	4000
2024-05-05 10:43:53	481657	car	4007
2024-05-05 10:43:54	2003656	car	4008
2024-05-05 10:43:56	5588831	truck	4000
2024-05-05 10:43:58	9615595	car	4004